

Architecture & case study

The problem

When a candidate is marked hired in Ashby, People Ops typically spends hours on manual work: notifying IT, creating the employee record in Rippling, writing a hiring manager brief, setting up 1:1s, adding the new hire to Slack channels, and preparing a first-week guide. Each step is a potential gap. Quality depends entirely on who is doing it and how much time they have.

The deeper problem: generic onboarding templates fail because they are built at the department level. A hiring manager brief for a Senior Engineer looks nothing like one for a Customer Success Manager. Building and maintaining role-specific templates for every position is unsustainable.

The architecture

A fully triggered, role-specific onboarding workflow that launches the moment a candidate is marked hired — requiring zero manual initiation from People Ops.

Step	System	What happens
1	Ashby	Candidate marked hired — webhook fires automatically
2	n8n	Trigger caught — pulls candidate name, role, department, employment type, start date, hiring manager, and job c
3	Claude API	JD passed as primary input — Claude reasons over the role and generates three outputs simultaneously
4a	Notion	Hiring manager brief created: pre-start prep, day 1 agenda, week 1 check-ins, 30/60/90 milestones
4b	Rippling	Employee record created + onboarding workflow tasks assigned to correct owners
4c	Jira	IT provisioning ticket created, pre-populated with role-specific access requirements
4d	Google Calendar	1:1 meetings scheduled between new hire and hiring manager for weeks 1-4
5	People Ops	Receives summary notification — reviews, approves, sends. No manual drafting.

Why the job description matters

The key architectural decision is passing the full job description to Claude as the primary input — not just the department or role title. This means every output is generated for that specific role, not a generic template. A hiring manager brief for a People Systems Specialist includes different tools, priorities, and stakeholders than one for an Account Executive. Claude reasons over the JD to determine what tools need provisioning, what the first week should focus on, and what the 30/60/90 milestones should look like.

The intelligence layer vs the automation layer

A common mistake is using AI for everything. This architecture uses AI only where variable, contextual judgment is required — and native automation everywhere else.

Claude API handles	Native automation handles
Hiring manager brief (role-specific)	Employee record creation in Rippling
New hire onboarding guide (role-specific)	IT ticket creation in Jira
IT access requirements (inferred from JD)	Calendar invite scheduling
First week priorities (inferred from JD)	Slack channel additions
30/60/90 milestones (inferred from JD)	Benefits enrollment triggers

What changes

Outcome area	Before → after
Onboarding prep time	Hours of manual work → minutes of review and approval
Output consistency	Depends on who is doing it → consistent regardless of workload
Hiring manager experience	Chase People Ops for a brief → receives one before asking
IT provisioning	Manual ticket with partial information → pre-populated ticket with role-specific access requirements
People Ops role	Executing onboarding → owning the system that executes it